



SARIMA (Seasonal Autoregressive Integrated Moving Average)

Last Updated : 24 May, 2024

Time series data is all around us, from stock prices and weather patterns to demand forecasting and seasonal trends in sales. To make sense of this data and predict future values, we turn to powerful models like the Seasonal Autoregressive Integrated Moving Average, or SARIMA. In this article, we will unravel the mysteries of SARIMA models, to forecast the monthly sales.

Understanding SARIMA

SARIMA, which stands for Seasonal Autoregressive Integrated Moving Average, is a versatile and widely used time series forecasting model. It's an extension of the non-seasonal ARIMA model, designed to handle data with seasonal patterns. SARIMA captures both short-term and long-term dependencies within the data, making it a robust tool for forecasting. It combines the concepts of autoregressive (AR), integrated (I), and moving average (MA) models with seasonal components.

The Components of SARIMA

To grasp SARIMA, let's break down its components:

- Seasonal Component:** The "S" in SARIMA represents seasonality, which refers to repeating patterns in the data. This could be daily, monthly, yearly, or any other regular interval. Identifying and modelling the seasonal component is a key strength of SARIMA.
- Autoregressive (AR) Component:** The "AR" in SARIMA signifies the autoregressive component, which models the relationship between the current data point and its past values. It captures the data's autocorrelation, meaning how correlated the data is with itself over time.
- Integrated (I) Component:** The "I" in SARIMA indicates differencing, which transforms non-stationary data into stationary data. Stationarity is crucial for time series modelling. The integrated component measures how many differences are required to achieve stationarity.
- Moving Average (MA) Component:** The "MA" in SARIMA represents the moving average component, which models the dependency between the current data point and past prediction errors. It helps capture short-term noise in the data.

Seasonal Differencing

Before we jump into SARIMA, it's essential to understand seasonal differencing. Seasonal differencing is the process of subtracting the time series data by a lag that equals the seasonality. This helps remove the seasonal component and makes the data stationary, allowing for more straightforward modeling. Seasonal differencing is often denoted as "D" in SARIMA.

The SARIMA Notation

SARIMA(p, d, q)(P, D, Q, s):

AR(p): Autoregressive component of order p

MA(q): Moving average component of order q

I(d): Integrated component of order d

Seasonal AR(P): Seasonal autoregressive component of order P

MA(Q): Seasonal moving average component of order Q

Seasonal I(D): Seasonal integrated component of order D

s: Seasonal period

The mathematical representation of SARIMA is as follows:

$$(1 - \phi_1 B)(1 - \Phi_1 B^s)(1 - B)(1 - B^s)y_t = (1 + \theta_1 B)(1 + \Theta_1 B^s)\varepsilon_t$$

Where,

- y_t is the observed time series at time t ,
- B is the backward shift operator, representing the lag operator (i.e., $By_t = y_{t-1}$)
- ϕ_1 is the non-seasonal autoregressive coefficient,
- Φ_1 is the seasonal autoregressive coefficient,
- θ_1 is the non-seasonal moving average coefficient,
- Θ_1 is the seasonal moving average coefficient,
- s is the seasonal period,
- ε_t is the white noise error term at time t .

Let's go through each component of the equation:

- **Autoregressive (AR) Component:** The autoregressive non-seasonal component represented by $(1 - \phi_1 B)$ captures the relationship between the current observation and a certain number of lagged observations (previous values in the time series). The B term represents the backshift operator is commonly used in time series analysis. It represents the lag operator, which shifts the time series backward by a certain number of time period. The order of the autoregressive component, denoted by (p) , determines the number of past values considered in the model.
- **Seasonal Autoregressive (SAR) Component:** The seasonal autoregressive component is represented by $(1 - \Phi_1 B^s)$. This component captures the relationship between the current observation and a certain number of lagged observations at seasonal intervals. The B^s term represents the backshift operator applied to the seasonal lagged observations.
- **Non-Seasonal Differencing Component:** The non-seasonal differencing component is represented by $(1 - B)$, where d is the order of non-seasonal differencing. This component is used to make the

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

- **Seasonal Differencing Component:** The seasonal differencing component is represented by $(1 - B^s)$, where D is the order of seasonal differencing and. This component is used to make the time series stationary by differencing it at seasonal intervals.
- **Observed Time Series:** The observed time series is denoted by y_t . It represents the historical data that we have and want to forecast.
- **Moving Average (MA) Component:** The moving average component is represented by $(1 + \theta_1 B)$. This component captures the relationship between the current observation and the residual errors from a moving average model applied to lagged observations.
- **Seasonal Moving Average (SMA) Component:** The seasonal moving average component is represented by $(1 + \Theta_1 B^s)$. This component captures the relationship between the current observation and the residual errors from a moving average model applied to lagged observations at seasonal intervals.
- **Error Term:** The error term is denoted by ε_t . It represents the random noise or unexplained variation in the time series.

The **SARIMA model** is estimated by fitting the model to the historical data and then using it to forecast future values.

Use Cases for SARIMA

SARIMA models find applications in various domains, including:

- **Economics:** Predicting economic indicators like inflation and GDP.
- **Retail:** Forecasting sales and demand for seasonal products.
- **Energy:** Predicting energy consumption and demand.
- **Healthcare:** Modeling patient admissions and disease outbreaks.
- **Finance:** Predicting stock prices and market trends.

Implementation of SARIMA for Time Series Forecasting:

Step 1: Import Libraries

First, import the necessary libraries for time series analysis.

Python



```

4 from statsmodels.tsa.statespace.sarimax import SARIMAX
5 from statsmodels.tsa.stattools import adfuller
6 from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
7 from sklearn.metrics import mean_absolute_error, mean_squared_error

```

Step 2: Load Dataset

let's use this time retail dataset of a global superstore for 4 years, to predict the superstore's monthly sales.

Python

```

1 df=pd.read_csv('Dataset- Superstore (2015-2018).csv')
2 sales_data=df[['Order Date', 'Sales']]
3 sales_data=pd.DataFrame(sales_data)
4 sales_data['Order Date']=pd.to_datetime(sales_data['Order Date'])
5 print(sales_data.head())

```

Output:

	Order Date	Sales
0	2016-11-08	261.9600
1	2016-11-08	731.9400
2	2016-06-12	14.6200
3	2015-10-11	957.5775
4	2015-10-11	22.3680

Step 3: Extracting monthly sales

Let's extract monthly sales, as day by day sales analysis won't be of any use.

Python

```

1 df1 = sales_data.set_index('Order Date')
2
3 monthly_sales = df1.resample('M').mean()
4 monthly_sales.head()

```

Output:

	Sales
Order Date	
2014-01-31	180.213861
2014-02-28	98.258522
2014-03-31	354.719803
2014-04-30	209.595148
2014-05-31	193.838418

Step 5: Plotting the Monhly Sales

Python


We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

```

3 plt.title("Monthly sales")
4 plt.xlabel("Date")
5 plt.ylabel("Sales")
6 plt.show()

```

Output:



From the visualization, we can analyze that the company has a seasonal pattern in sales. The company has a positive correlation between its sales and the month of the year and a negative correlation between its sales and the month of the quarter.

Step 6: Check Stationarity

Before applying SARIMA, check if your time series data is stationary because SARIMA assumes that the time series data is stationary. Stationarity refers to the statistical properties of a time series remaining constant over time, such as constant mean, constant variance, and constant autocovariance. You can use the Dickey-Fuller test for this.

Perform the ADF Test: Inside the function, it calls the `adfuller` function on the timeseries. The `autolag='AIC'` parameter specifies that the lag order should be chosen based on the Akaike Information Criterion (AIC).

Retrieve the p-Value: It extracts the p-value from the ADF test result, which is stored in the variable `p_value`.

Print the Results: It prints the ADF Statistic (`result[0]`), the p-value (`p_value`), and a statement indicating whether the time series is stationary or non-stationary based on the p-value. If the p-value is less than 0.05, it considers the series as "Stationary"; otherwise, it's labeled as "Non-Stationary."

Check Stationarity of sales: The function is then called with the sales time series data, which is assumed to be the monthly sales data.

Python

```

1 def check_stationarity(timeseries):
2     # Perform the Dickey-Fuller test
3     result = adfuller(timeseries, autolag='AIC')
4     p_value = result[1]
5     print(f'ADF Statistic: {result[0]}')
6     print(f'p-value: {p_value}')
7     print('Stationary' if p_value < 0.05 else 'Non-Stationary')
8
9
10 check_stationarity(monthly_sales['Sales'])

```

Output:

```
ADF Statistic: -3.2865668298704227
```

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

Now, as the data is stationary, we can proceed further to build a model.

Step 7: Identify Model Parameters

Identify the SARIMA model parameters (p, d, q, P, D, Q, s) using ACF and PACF plots.


ACF Plot: This function generates an ACF plot, which is a plot of autocorrelations of the differenced time series. Autocorrelation measures the relationship between a data point and previous data points at different lags.

PACF Plot: This function generates a PACF plot, which is a plot of partial autocorrelations of the differenced time series. Partial autocorrelation represents the correlation between a data point and a lag while adjusting for the influence of other lags.

Python

```
1 # Plot ACF and PACF
2 plot_acf(monthly_sales)
3 plot_pacf(monthly_sales)
4 plt.show()
```

Output:



We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

Here, both plots have a horizontal axis that shows the lags, and a vertical axis representing correlation coefficients ranging from -1 to 1 where -1 means perfect negative correlation, 0 means no correlation and 1 means perfect positive correlation.

The blue shaded area represents confidence interval for the correlation coefficients. **If the correlation coefficient at a certain lag is outside the confidence interval, it means that the correlation coefficient is statistically significant and not due to chance.**

Step 8: Fit the SARIMA Model

Now, fit the SARIMA model using the identified parameters.

SARIMA Parameters: These parameters define the SARIMA model. p, d, and q represent the non-seasonal order, and P, D, Q, and s represent the seasonal order. For example, p is the autoregressive (AR) order, d is the differencing order, and s is the seasonality (in this case, 12 for monthly seasonality).

Fitting the SARIMA Model: Here, a SARIMAX model is created using the defined parameters. The SARIMAX function takes the sales data and the specified order and seasonal_order parameters. The fit method is then called to estimate the model's coefficients based on the data.

Python

```
1 # Define SARIMA parameters
2 p, d, q = 1, 1, 1
3 P, D, Q, s = 1, 1, 1, 12 # Assuming monthly seasonality
4
5 # Fit the SARIMA model
6 model = SARIMAX(monthly_sales, order=(p, d, q), seasonal_order=(P, D, Q, s))
7 results = model.fit()
8 model
```

Output:

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 5 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 4.15152D+00 |proj g|= 2.74080D-01

At iterate 5 f= 4.06571D+00 |proj g|= 1.57031D-02

At iterate 10 f= 4.05982D+00 |proj g|= 4.46554D-03

At iterate 15 f= 4.03337D+00 |proj g|= 1.65496D-02

At iterate 20 f= 4.02336D+00 |proj g|= 3.56570D-02

At iterate 25 f= 4.01887D+00 |proj g|= 6.31166D-04

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
5	28	36	1	0	0	5.795D-06	4.019D+00

F = 4.0188723995389912

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL

Step 9: Generate Forecasts

With the model fitted, generate forecasts for future time periods.

In this section of the code, we are using the SARIMA model to forecast future sales values. We specify that we want to forecast the next 12 months by setting forecast_periods to 12.

The results.get_forecast(steps=forecast_periods) function generates the forecast, providing both the predicted mean and a confidence interval for the forecasted values.

Python

```
1 # Forecast future values
2 forecast_periods = 12 # Forecast the next 12 months
3 forecast = results.get_forecast(steps=forecast_periods)
4 forecast_mean = forecast.predicted_mean
5 forecast_ci = forecast.conf_int()
6
```


We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

```

10 plt.plot(forecast_mean, label='Forecast', color='red')
11 plt.fill_between(forecast_ci.index, forecast_ci.iloc[:, 0], forecast_ci.iloc[:, 1], color='pink')
12 plt.title("Sales Forecast")
13 plt.xlabel("Date")
14 plt.ylabel("Sales")
15 plt.legend()
16 plt.show()

```

Output:



The observed sales data and the forecasted values in red. The pink shaded area represents the confidence interval around the forecast. This visualization helps us understand the expected future sales trends based on our SARIMA model.

Step 10: Evaluate the Model

Let's evaluate the forecasted sales values by comparing them to the observed sales data using two common metrics for this evaluation: Mean Absolute Error (MAE) and Mean Squared Error (MSE).

MAE (Mean Absolute Error) measures the average absolute difference between the observed and forecasted values. It provides a simple and easily interpretable measure of the model's accuracy.

MSE (Mean Squared Error) measures the average of the squared differences between the observed and forecasted values. MSE gives more weight to large errors and is sensitive to outliers.

Lower values indicate better performance.

Python

```

1 observed = monthly_sales[-forecast_periods:]
2 mae = mean_absolute_error(observed, forecast_mean)
3 mse = mean_squared_error(observed, forecast_mean)
4 print(f'MAE: {mae}')
5 print(f'MSE: {mse}')

```

Output:

MAE: 30.815697783063925

MSE: 1302.515608316509

Conclusion

SARIMA models are a potent tool in the realm of time series forecasting. They can capture complex seasonality, long-term trends, and short-term fluctuations in your data. By understanding the components, notation, and best practices of SARIMA modeling, you can make accurate predictions in various domains. So, the next time you encounter time series data, remember the power of SARIMA to unravel its hidden patterns and forecast the future with confidence.

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

practices, hands-on projects, and real-world applications. Become the executive head of industries related to **Data Analysis**, **Machine Learning**, and **Data Visualization** with these growing skills. Ready to Transform Your Future? **Enroll Now to Be a Data Science Expert!**

[Comment](#)[More info](#)[Next Article](#)

Similar Reads

SARIMA (Seasonal Autoregressive Integrated Moving Average)

Time series data is all around us, from stock prices and weather patterns to demand forecasting and seasonal trends in sales. To make sense of this data and predict future values, we turn to powerful models like the Seasonal Autoregressive Integrated Moving Average, or SARIMA. In this article, we wi

11 min read

Vector Autoregressive Model (VAR) Using R

Vector Autoregressive (VAR) models are a fundamental tool in time series analysis, particularly useful for multivariate time series data. VAR models capture the linear interdependencies among multiple time series. This article will provide a comprehensive overview of VAR models and demonstrate...

6 min read

Autoregressive (AR) Model for Time Series Forecasting

Autoregressive models, often abbreviated as AR models, are a fundamental concept in time series analysis and forecasting. They have widespread applications in various fields, including finance, economics, climate science, and more. In this comprehensive guide, we will explore autoregressive...

11 min read

Vector Autoregression (VAR) for Multivariate Time Series

Vector Autoregression (VAR) is a statistical tool used to investigate the dynamic relationships between multiple time series variables. Unlike univariate autoregressive models, which only forecast a single variable based on its previous values, VAR models investigate the interconnectivity of many va

6 min read

Write a brief note on Inter Tropical Convergence Zone?

During the monsoon season, India's climate is hot and muggy, similar to that of South and Southeast Asia. Two of the four seasons that exist in the Indian subcontinent are monsoon seasons. Those are Season of the Southwest Monsoon and The Monsoon Season in the North-East. The Arabic term mausin...

5 min read

Compute a Moving Average in MySQL

Moving the averages is an important technique in the field of data analysis and time-series data processing. By reducing data fluctuations moving averages aid in highlighting trends in the data. We'll examine how to compute moving averages in MySQL a well-liked relational database administratio...

3 min read

Moving Averages in R

Moving averages (MA) play a crucial role in time series analysis, providing a means to discern trends, patterns, and underlying structures within sequential data. In the context of R, a popular programming language for statistical computing, various functions and packages facilitate the computation

6 min read

Add Moving Average Plot to Time Series Plot in R

In time series analysis, a moving average is a widely used technique to smooth out short-term fluctuations and highlight longer-term trends or cycles. R

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

How to calculate MOVING AVERAGE in a Pandas DataFrame?

Calculating the moving average in a Pandas DataFrame is used for smoothing time series data and identifying trends. The moving average, also known as the rolling mean, helps reduce noise and highlight significant patterns by averaging data points over a specific window. In Pandas, this can be achieved

7 min read

Compute Moving Average in PostgreSQL

PostgreSQL is an advanced relational database system that supports both relational (SQL) and non-relational (JSON) queries. It is free and open-source. The moving average helps to level the price data over a specified period by creating a constantly updated average price. PostgreSQL which, is an imp

6 min read

Understanding the Moving average (MA) in Time Series Data

Data is often collected with respect to time, whether for scientific or financial purposes. When data is collected in a chronological order, it is referred to as time series data. Analyzing time series data provides insights into how the data behaves over time, including underlying patterns that can

15 min read

How to Decide Window Size for a Moving Average Filter in MATLAB?

A moving average filter is a widely used technique for smoothing data in signal processing. It is used to reduce the amount of noise in a given signal and to identify trends in the data. In MATLAB, the window size of a moving average filter is an important parameter that determines how much data is

3 min read

Seasonal Rhythm| Class 11 Geography Notes

Seasonal Rhythm is a subpart of the Class 11 Geography Chapter "Climate and Vegetation". The Seasonal Rhythm tells us about the various occurrences related to the regular and cyclic pattern of the different changes in the weather, temperature, and different flora and fauna. Earth's position and its

5 min read

Seasonal Adjustment and Differencing in Time Series

Time series data can be difficult to evaluate successfully because of the patterns and trends it frequently displays. To address these tendencies and improve the data's suitability for modeling and analysis, two strategies are employed: seasonal adjustment and differencing. Table of Content Seasonal

11 min read

How to Compute a Moving Average in PL/SQL?

A moving average is a technique we use to analyze and determine some specific trends in our provided data. Through moving averages, we can analyze and determine trends in our data along with some other benefits like noise reduction in our data. In this article, we are going to learn about "how to co

5 min read

Complete Guide To SARIMAX in Python

Time series data consists of observations collected over time at equally spaced intervals. SARIMAX is a statistical model designed to capture and forecast the underlying patterns, trends, and seasonality in such data. In this article, we'll explore the SARIMAX model, understand its mathematical unde

12 min read

Time Series Data Transformation Using R

Time series data transformation is important to prepare the data for modeling and making informed and correct predictions. It involves several steps that prepare the data such that it doesn't give wrong predictions. R is a statistical programming language that is widely used by researchers and analy

7 min read

How to Calculate Moving Averages in Python?

In this discussion we are going to see how to Calculate Moving Averages in Python in this discussion we will write a proper explanation What is Moving Averages? Moving Averages, a statistical method in data analysis, smooths fluctuations in time-series data to reveal underlying trends. Calculating th

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

In this article, we will look the how to Calculate an Exponential Moving Average in R Programming Language. Exponential moving average (EMA) tells us the weighted mean of the previous K data points. EMA places a greater weight and significance on the most recent data points. To get the EMA we will u

2 min read

Article Tags :

[Machine Learning](#)

[Geeks Premier League](#)

[AI-ML-DS](#)

[Machine Learning](#)

Practice Tags :

[Machine Learning](#)

[Machine Learning](#)



Corporate & Communications Address:-
A-143, 7th Floor, Sovereign Corporate
Tower, Sector- 136, Noida, Uttar Pradesh
(201305) | Registered Address:- K 061,
Tower K, Gulshan Vivante Apartment,
Sector 137, Noida, Gautam Buddh
Nagar, Uttar Pradesh, 201305



Company

[About Us](#)
[Legal](#)
[In Media](#)
[Contact Us](#)
[Advertise with us](#)
[GFG Corporate Solution](#)
[Placement Training Program](#)
[GeeksforGeeks Community](#)

DSA

[Data Structures](#)
[Algorithms](#)
[DSA for Beginners](#)
[Basic DSA Problems](#)
[DSA Roadmap](#)
[Top 100 DSA Interview Problems](#)
[DSA Roadmap by Sandeep Jain](#)
[All Cheat Sheets](#)

Web Technologies

[HTML](#)
[CSS](#)
[JavaScript](#)
[TypeScript](#)

Languages

[Python](#)
[Java](#)
[C++](#)
[PHP](#)
[GoLang](#)
[SQL](#)
[R Language](#)
[Android Tutorial](#)
[Tutorials Archive](#)

Data Science & ML

[Data Science With Python](#)
[Data Science For Beginner](#)
[Machine Learning](#)
[ML Maths](#)
[Data Visualisation](#)
[Pandas](#)
[NumPy](#)
[NLP](#)
[Deep Learning](#)

Python Tutorial

[Python Programming Examples](#)
[Python Projects](#)
[Python Tkinter](#)
[Web Scraping](#)

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

Web Design

Computer Science

Operating Systems
Computer Network
Database Management System
Software Engineering
Digital Logic Design
Engineering Maths
Software Development
Software Testing

System Design

High Level Design
Low Level Design
UML Diagrams
Interview Guide
Design Patterns
OOAD
System Design Bootcamp
Interview Questions

School Subjects

Mathematics
Physics
Chemistry
Biology
Social Science
English Grammar
Commerce
World GK

DevOps

Git
Linux
AWS
Docker
Kubernetes
Azure
GCP
DevOps Roadmap

Interview Preparation

Competitive Programming
Top DS or Algo for CP
Company-Wise Recruitment Process
Company-Wise Preparation
Aptitude Preparation
Puzzles

GeeksforGeeks Videos

DSA
Python
Java
C++
Web Development
Data Science
CS Subjects

@GeeksforGeeks, Sanchhaya Education Private Limited, All rights reserved